

Title: Predicting Employee Salary Based on Experience

Problem Statement:

Background:

In the corporate world, employee compensation is a crucial factor for both the employers and the employees. Determining a fair and competitive salary based on an employee's experience is important for maintaining job satisfaction, motivation, and retention. This dataset contains data on employees' years of experience and their corresponding salaries.

Objective:

The objective of this analysis is to build a predictive model that can accurately forecast an employee's salary based on their years of experience. This model will help in understanding the salary trends related to experience and assist companies in establishing fair compensation practices.

Dataset Description:

The dataset consists of the following columns: 1.Experience_Years: Number of years of experience the employee has. 2.Salary: Salary of the employee (in dollars).

Importing Libraries

```
In [53]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Import DataSet

```
In [2]: emp_sal=pd.read_csv("salary_exp.csv")
emp_sal
```

Out[2]:

	Experience Years	Salary
0	1.1	39343
1	1.2	42774
2	1.3	46205
3	1.5	37731
4	2.0	43525
5	2.2	39891
6	2.5	48266
7	2.9	56642
8	3.0	60150
9	3.2	54445
10	3.2	64445
11	3.5	60000
12	3.7	57189
13	3.8	60200
14	3.9	63218
15	4.0	55794
16	4.0	56957
17	4.1	57081
18	4.3	59095
19	4.5	61111
20	4.7	64500
21	4.9	67938
22	5.1	66029
23	5.3	83088
24	5.5	82200
25	5.9	81363
26	6.0	93940
27	6.2	91000
28	6.5	90000
29	6.8	91738
30	7.1	98273
31	7.9	101302
32	8.2	113812
33	8.5	111620
34	8.7	109431
35	9.0	105582
36	9.5	116969
37	9.6	112635
38	10.3	122391
39	10.5	121872

Data Understanding

In [3]:

```
emp_sal.head()
```

```
Out[3]:
```

	Experience Years	Salary
0	1.1	39343
1	1.2	42774
2	1.3	46205
3	1.5	37731
4	2.0	43525

Initial Check Up

```
In [4]: emp_sal.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 40 entries, 0 to 39
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Experience Years  40 non-null     float64
1   Salary           40 non-null     int64  
dtypes: float64(1), int64(1)
memory usage: 772.0 bytes

In [6]: emp_sal.describe()
```

```
Out[6]:
```

	Experience Years	Salary
count	40.000000	40.000000
mean	5.152500	74743.625000
std	2.663715	25947.122885
min	1.100000	37731.000000
25%	3.200000	56878.250000
50%	4.600000	64472.500000
75%	6.875000	95023.250000
max	10.500000	122391.000000

```
In [7]: emp_sal.shape

Out[7]: (40, 2)
```

Asking Questions to the Data


1.What is the highest Salary of employee and how many years of experience employee has. 2.Years of experience greater than 7. 3.Which Experience year has Salary is equals to 46205. 4.Maximum salary of an employee 5.Average salary of an employee 6.Maximum years of experience 7.Average year experience of an employee 8.How many null values are in the emp_sal

Data Visualization

```
In [8]: #1.What is the highest Salary of employee and how many years of experience employee has.
#4.Maximun salary of an employee
#7.Average years of experience

sns.scatterplot(emp_sal,x="Experience Years",y="Salary")

Out[8]: <Axes: xlabel='Experience Years', ylabel='Salary'>
```



```
In [49]: #2.Years of experience greater than 7.
emp_sal[emp_sal['Experience Years']>7]
```

```
Out[49]:
```

	Experience Years	Salary
30	7.1	98273
31	7.9	101302
32	8.2	113812
33	8.5	111620
34	8.7	109431
35	9.0	105582
36	9.5	116969
37	9.6	112635
38	10.3	122391
39	10.5	121872

```
In [50]: #3.Which Experience year has Salary is equals to 46205.
emp_sal[emp_sal['Salary']==46205]
```


```
Out[50]:
```

	Experience Years	Salary
2	1.3	46205

```
In [11]: #5. Average salary of an employee
sns.kdeplot(emp_sal,x="Salary")
```

C:\Users\sindhu\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):


```
Out[11]: <Axes: xlabel='Salary', ylabel='Density'>
```



```
In [12]: #7.Average experience of an employee
sns.histplot(emp_sal,x='Experience Years')
```

C:\Users\sindhu\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

```
Out[12]: <Axes: xlabel='Experience Years', ylabel='Count'>
```




```
In [13]: #8.How many null values are in the emp_sal
emp_sal.isnull().count()
```

```
Out[13]: Experience Years    40
Salary                    40
dtype: int64
```

```
In [14]: sns.kdeplot(data=emp_sal,x='Experience Years')
```

C:\Users\sindhu\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):


```
Out[14]: <Axes: xlabel='Experience Years', ylabel='Density'>
```



```
In [15]: sns.pairplot(emp_sal)
```

C:\Users\sindhu\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
 with pd.option_context('mode.use_inf_as_na', True):
 C:\Users\sindhu\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
 with pd.option_context('mode.use_inf_as_na', True):

```
Out[15]: <seaborn.axisgrid.PairGrid at 0x2f5f8476b10>
```



ML

Linear Regression

Linear regression is a supervised machine learning method that provides a linear relationship between an independent variable and a dependent variable to predict the outcome of future events.

```
In [17]: # correlation
```

```
emp_sal.corr()
```

```
Out[17]:
```

	Experience Years	Salary
Experience Years	1.000000	0.977692
Salary	0.977692	1.000000

```
In [18]: emp_sal.max()
```

```
Out[18]: Experience Years      10.5  
Salary      122391.0  
dtype: float64
```

```
In [19]: emp_sal.min()
```

```
Out[19]: Experience Years      1.1  
Salary      37731.0  
dtype: float64
```

```
In [20]: X=emp_sal.iloc[:,0:1]  
y=emp_sal.iloc[:,-1]  
X.head()
```

```
Out[20]:
```

	Experience Years
0	1.1
1	1.2
2	1.3
3	1.5
4	2.0

```
In [21]: y.head()
```

```
Out[21]: 0    39343  
1    42774  
2    46205  
3    37731  
4    43525  
Name: Salary, dtype: int64
```

Train and Test

```
In [ ]:
```

```
In [22]: from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=2)
```

```
In [23]: X_train.head()
```

```
Out[23]:
```

	Experience Years
17	4.1
37	9.6
38	10.3
29	6.8
24	5.5

```
In [24]: X_test
```

```
Out[24]:
```

	Experience Years
27	6.2
9	3.2
14	3.9
0	1.1
2	1.3
30	7.1
13	3.8
36	9.5

```
In [25]: y_test
```

```
Out[25]: 27    91000
          9    54445
          14   63218
          0    39343
          2    46205
          30   98273
          13   60200
          36  116969
          Name: Salary, dtype: int64
```

```
In [26]: X_train.shape
```

```
Out[26]: (32, 1)
```

```
In [27]: X_test.shape
```

```
Out[27]: (8, 1)
```

```
In [28]: y_train.shape
```

```
Out[28]: (32,)
```

```
In [29]: y_test.shape
```

```
Out[29]: (8,)
```

Model Building

```
In [30]: from sklearn.linear_model import LinearRegression
         Lr=LinearRegression()
         Lr
```


```
Out[30]: ▼ LinearRegression
         LinearRegression()
```

```
In [31]: Lr.fit(X_train,y_train)
```

```
Out[31]: ▼ LinearRegression
         LinearRegression()
```

```
In [32]: #Data visualization
         import matplotlib.pyplot as plt
         plt.scatter(emp_sal['Experience Years'],emp_sal['Salary'])
         plt.plot(X_train,Lr.predict(X_train),color='red')
         plt.xlabel("Experience Years")
         plt.ylabel("Salary")
```

```
Out[32]: Text(0, 0.5, 'Salary')
```



#calculating coefficient, Intercept and predicted value


```

In [33]: Lr.coef_

Out[33]: array([9629.89561636])

In [34]: Lr.intercept_

Out[34]: 24469.054538114055

In [35]: #predict
Lr.predict(X_test)

Out[35]: array([ 84174.40735952,  55284.72051045,  62025.6474419 ,  35061.9397161 ,
                36987.91883938,  92841.31341423,  61062.65788026, 115953.06289349])

In [54]: Lr.predict([[3.2]])

C:\Users\sindhu\anaconda3\Lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature name
s, but LinearRegression was fitted with feature names
warnings.warn(

Out[54]: array([55284.72051045])

```

The predicted salary of 3.2 is 55284 and the actual value of 3.2 is 64445. The predicted value is decreased by 9161 than actual value.

Model Evaluation

```

In [37]: from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
y_pred=Lr.predict(X_test)
y_pred

Out[37]: array([ 84174.40735952,  55284.72051045,  62025.6474419 ,  35061.9397161 ,
                36987.91883938,  92841.31341423,  61062.65788026, 115953.06289349])

In [38]: mean_absolute_error(y_test,y_pred)

Out[38]: 3708.261090762284

In [39]: r2_score(y_test,y_pred)

Out[39]: 0.9655807830897453

In [40]: MSE=mean_squared_error(y_test,y_pred)
MSE

Out[40]: 22909642.289620496

In [41]: import pandas as pd
from sklearn.metrics import mean_squared_error
RMSE=MSE**0.5
data_rmse={'Actual (y_test)':y_test,'predicted(y_pred)':y_pred}
df_rmse=pd.DataFrame(data_rmse)
df_rmse.head()

Out[41]:
   Actual (y_test)  predicted(y_pred)
27             91000             84174.407360
9              54445             55284.720510
14             63218             62025.647442
0              39343             35061.939716
2              46205             36987.918839

In [43]: df_rmse['Actual (y_test)'].sum()

Out[43]: 569653

In [42]: # Total Prediction
df_rmse['predicted(y_pred)'].sum()

Out[42]: 543391.668055328

```

Insights

```

In [2]: #1.In salary_exp dataset consists of two columns "Experience Years" and "salary".
#2.Salary - Represents the salary of an employee.
#3.Experience Years - Represents the years of experience.
#4.'salary' dtype - int, 'Experience Years' dtype - float and Both columns has 40 null values.
#5.Average years of experience is 5.1 years of an employee.

```

#6.Average amount of salary that employee has \$74000.
#7.Highest salary of an employee is \$122000.
#8.Maximum years of experience an employee has 10.5
#9.Minimum salary is \$37000 and minimum Experience Years is 1.1 years.
#10.Shape of emp_sal dataset is(40,2) means 40 rows and 2 columns.
#11.Shape of X_test,y_test is(8,1) and X_train,y_train is(32,1)
#12.iloc is used for integer indexing values 'X' has Experience Years & 'y' has salary column.
#13.corr() tells the relationship between Salary and Experience Years.
#14.fit() can be used to fit data into LinearRegression.
#15.LinearRegression helps to determine the relationship between Experience Years and salary.
#16.LinearRegression allows for prediction of salary based on given years of experience by using predict().
#17.Scatter plot to visualize the relationship between Experience Years and salary.
#18.kde plot for both Experience Years and salary to understand data visualization.
#19.Calculate Model Evaluation metrics such as MAE,MSE and R2-squared helps in assess model's performance.
#20.The predicted value of 3.2 is decreased by 9161 than actual value of 3.2.